

IT IS SLOW

...

performance analysis and optimisation on the JVM

About

Galo Navarro ~ @srvaroa

- Software engineer: backend, distributed systems, data plumbing
- [...], Last.fm (2008), Tuenti (2011), Midokura (2013), Zhilabs (2016)

Agenda

- Motivation & goals
- Getting answers to:
 - Do I have performance problems?
 - Which one?
 - How can I fix them?
- Takeaways
 - Performance matters
 - Develop with performance in mind
 - Toolbox

Why should I care about performance?

Performance limits business goals

- Resiliency → is your business functioning?
 - DoS: Denial Of Service (or, Die Of Success: Hacker News / Digg / Reddit effect)
- Efficiency
 - Throwing hardware at the problem is not a silver bullet
 - You don't want "distributed systems" in your problem set
 - Cost is a factor for potential customers (cautionary tale: Scylla DB vs. Cassandra)
 - Not an option in ARM, mobile devices

Why should I care about performance?

Are performance metrics explicit in your business requirements?

What are your SLAs?

How do latency, throughput.. relate to your business targets?

- Amazon: "it is estimated that a 100-millisecond delay reduces Amazon's sales by 1 percent." [1]
- Google: “..half a second delay caused a 20% drop in traffic”
- Yours?

[1]: <http://www.nytimes.com/2009/06/14/magazine/14search-t.html>

[2]: <http://glinden.blogspot.com.es/2006/11/marissa-mayer-at-web-20.html>

[3]: <http://perspectives.mvdirona.com/2009/10/the-cost-of-latency/>

Performance analysis.. on the JVM


Brendan Gregg: “Systems Performance” (2013). Highlights need for:

- Methodologies (e.g.: USE, TSA..)
- Checklists
- Toolbox

“Linux Performance Analysis in 60.000 ms”

- Protocol for first approach to performance-related incident

Proposes systematising specific approaches for common services (MySQL, Cassandra, Apache..), VMs (Java, Go, Ruby..), etc.



Focusing on low level tools

Profilers

- VisualVM, YourKit, etc. (use instrumentation & JVM Tool Interface)
- Flight Recorder + Mission Control (better: use internal JVM counters & APIs)

Caveats:

- Not always usable on production servers
 - Customers often deny access to their environment
- Focused on dev, forensics, not during incidences or downtime
- Licensing, vendor lock-in (e.g.: Flight Recorder)

MEMORY

...

Errors

Common originators of an investigation

- Out of Memory

- `java.lang.OutOfMemoryError: Java heap space`

~8 possible reasons



- Too much GC

- `java.lang.OutOfMemoryError: GC overhead limit exceeded`

- OOM killer (Linux)

- `$ dmesg | grep "Out of memory"`

- `kernel: Out of memory: Kill process 746 (..) score 1822 or sacrifice child`

- Code Cache

- `VM warning: CodeCache is full. Compiler has been disabled.`

- Allocation/Promotion failure, to-space exhausted..

- this one is fine: collection required to make room in Eden, Survivor, Old Gen..

Memory footprint

Real memory may be (much) bigger than set by Xmx

```
top - 15:45:17 up 154 days, 16:16, 4 users, load average: 0.20, 0.69, 0.82
Tasks: 416 total, 1 running, 415 sleeping, 0 stopped, 0 zombie
Cpu(s): 0.2%us, 0.1%sy, 0.0%ni, 99.7%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 258313M total, 223978M used, 34334M free, 928M buffers
Swap: 268140M total, 1057M used, 267083M free, 47252M cached

  PID USER      PR  NI  VIRT  RES  SHR S %CPU  %MEM    TIME+  COMMAND
 53071 tomcat    20   0 143g 141g 636 S   2 56.0 1029:56 /usr/java/jre1.8.0_31/bin/java -server -XX:+UseG1GC -XX:SurvivorRatio=2 -Xmx100g
41123 tomcat    20   0 33.2g 24g 14m S   0 9.7 84:34.44 /etc/alternatives/jre/bin/java -Dmail.smtp.ehlo=false -Dmail.smtp.auth=false -c1
```

- top / htop expose real mem usage
- Off-heap:
 - Code cache Tune with InitialCodeCacheSize / ReservedCodeCacheSize
 - Off-heap buffers Tune with -XX:MaxDirectMemorySize=64M
 - Thread stacks Tune with -Xss=1024k
- Reclamation: note that batch jobs may hoard memory between runs

Utilisation & Saturation (GC)

```
$ jstat -options
```

```
-class
```

```
-compiler
```

```
-gc
```

```
-gccapacity
```

```
-gccause
```

```
-gcmetacapacity
```

```
-gcold
```

```
-gcoldcapacity
```

```
-gcold
```

```
-gcoldcapacity
```

```
-gcutil
```

```
-printcompilation
```

```
jstat -gc 3441 500
```

S0C	S1C	S0U	S1U	EC	EU	OC	OU	MC	MU	CCSC	CCSU	YGC	YGCT	FGC	FGCT	GCT
512,0	512,0	32,0	0,0	1705984,0	204863,2	412160,0	294957,4	21552,0	12570,9	4144,0	1122,7	137682	332,212	2	0,159	332,371
512,0	512,0	32,0	0,0	1960448,0	117701,7	412160,0	294957,4	21552,0	12570,9	4144,0	1122,7	137684	332,217	2	0,159	332,376
512,0	512,0	32,0	0,0	2253312,0	0,0	412160,0	294957,4	21552,0	12570,9	4144,0	1122,7	137686	332,223	2	0,159	332,382
512,0	512,0	0,0	32,0	2157568,0	1294911,1	412160,0	294957,4	21552,0	12570,9	4144,0	1122,7	137687	332,225	2	0,159	332,384
512,0	512,0	0,0	32,0	1979904,0	752715,0	412160,0	294957,4	21552,0	12570,9	4144,0	1122,7	137689	332,230	2	0,159	332,390
512,0	512,0	0,0	32,0	1819136,0	327692,9	412160,0	294957,4	21552,0	12570,9	4144,0	1122,7	137691	332,235	2	0,159	332,394
512,0	512,0	0,0	32,0	1673216,0	0,0	412160,0	294957,4	21552,0	12570,9	4144,0	1122,7	137693	332,239	2	0,159	332,399
512,0	512,0	0,0	32,0	1541120,0	215881,7	412160,0	294957,4	21552,0	12570,9	4144,0	1122,7	137695	332,244	2	0,159	332,404
512,0	512,0	0,0	32,0	2216960,0	443453,5	412160,0	294957,4	21552,0	12570,9	4144,0	1122,7	137697	332,250	2	0,159	332,409
512,0	512,0	0,0	32,0	2033152,0	0,0	412160,0	294957,4	21552,0	12570,9	4144,0	1122,7	137699	332,254	2	0,159	332,414
512,0	512,0	32,0	0,0	1948160,0	1442458,9	412160,0	294957,4	21552,0	12570,9	4144,0	1122,7	137700	332,257	2	0,159	332,417
512,0	512,0	32,0	0,0	1789952,0	1253723,3	412160,0	294957,4	21552,0	12570,9	4144,0	1122,7	137702	332,262	2	0,159	332,421
512,0	512,0	32,0	0,0	1647104,0	1351500,6	412160,0	294957,4	21552,0	12570,9	4144,0	1122,7	137704	332,267	2	0,159	332,426
512,0	512,0	0,0	32,0	1817600,0	0,0	412160,0	294957,4	21552,0	12570,9	4144,0	1122,7	137706	332,272	2	0,159	332,434
512,0	512,0	0,0	32,0	1671680,0	0,0	412160,0	294957,4	21552,0	12570,9	4144,0	1122,7	137708	332,277	2	0,159	332,438
512,0	512,0	0,0	32,0	1539584,0	339015,7	412160,0	294957,4	21552,0	12570,9	4144,0	1122,7	137710	332,282	2	0,159	332,442

Lots of garbage, but
no promotions

XC = capacity (KiB)
XU = utilisation (KiB)

Utilisation & Saturation (GC)


Understanding application pauses and latency spikes

```
-Xloggc:$PATH                                consider ramdisk / ssd [1]
-XX:+PrintGCDetails
-XX:+PrintClassHistogram
-XX:+PrintTenuringDistribution
-XX:+PrintPromotionFailure
-XX:+PrintGCApplicationStoppedTime
-XX:+UseGCLogFileRotation
-XX:NumberOfGCLogFiles=$NUM_FILES             default 1
-XX:GCLogFileSize=$SIZE[M|K]                  default 512k
-XX:+PrintAdaptiveSizePolicy
```

[1]: <https://engineering.linkedin.com/blog/2016/02/eliminating-large-jvm-gc-pauses-caused-by-background-io-traffic>

“The app is frozen for 2 seconds..”

Happy case

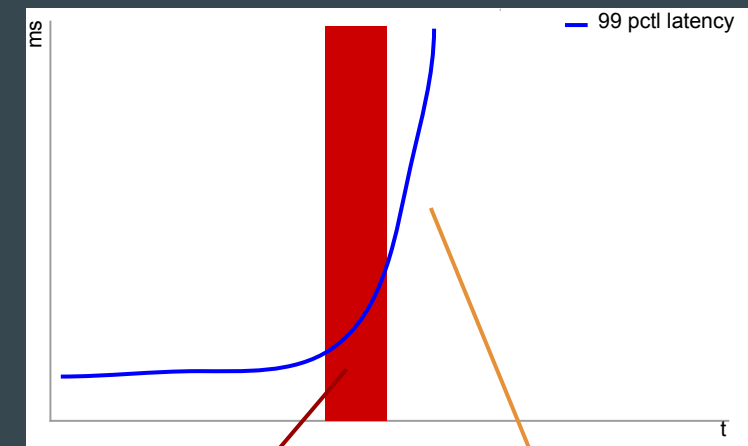


The event itself
requests pile up..

After recovery, system
deals with the queue

“The app is frozen for 2 seconds..”


Unhappy case: collapse



the event itself

too much backlog, overload

Middle ground: graceful degradation



survive with backpressure

Either way, this service's clients are not happy

2016-02-25T11:58:21.628+0800: [GC pause (G1 Evacuation Pause) (young)]

Desired survivor size 8053063680 bytes, new threshold 4 (max 15)

- age 1: 609830392 bytes, 609830392 total
- age 2: 635249376 bytes, 1245079768 total
- age 3: 530928792 bytes, 1776008560 total
- age 4: 6566883776 bytes, 8342892336 total
- age 5: 160917504 bytes, 8503809840 total
, 2.3754150 secs]

[Parallel Time: 2305.9 ms, GC Workers: 23]

[GC Worker Start (ms): Min: 63546061.8, Avg: 63546062.1, Max: 63546062.4, Diff: 0.6]

[Ext Root Scanning (ms): Min: 0.0, Avg: 0.3, Max: 0.6, Diff: 0.5, Sum: 6.3]

[SATB Filtering (ms): Min: 0.0, Avg: 0.0, Max: 0.1, Diff: 0.1, Sum: 0.1]

[Update RS (ms): Min: 16.7, Avg: 18.3, Max: 22.8, Diff: 6.1, Sum: 420.7]

[Processed Buffers: Min: 5, Avg: 10.4, Max: 18, Diff: 13, Sum: 239]

[Scan RS (ms): Min: 247.8, Avg: 252.3, Max: 253.9, Diff: 6.1, Sum: 5803.5]

[Code Root Scanning (ms): Min: 0.0, Avg: 0.0, Max: 0.0, Diff: 0.0, Sum: 0.3]

[Object Copy (ms): Min: 2032.6, Avg: 2033.1, Max: 2034.4, Diff: 1.8, Sum: 46762.1]

[Termination (ms): Min: 0.0, Avg: 1.3, Max: 1.7, Diff: 1.7, Sum: 30.3]

[GC Worker Other (ms): Min: 0.0, Avg: 0.1, Max: 0.2, Diff: 0.2, Sum: 2.1]

[GC Worker Total (ms): Min: 2305.2, Avg: 2305.5, Max: 2305.8, Diff: 0.6, Sum: 53025.4]

[GC Worker End (ms): Min: 63548367.5, Avg: 63548367.6, Max: 63548367.7, Diff: 0.2]

[Code Root Fixup: 0.1 ms]

[Code Root Migration: 0.1 ms]

[Code Root Purge: 0.0 ms]

[Clear CT: 2.1 ms]

[Other: 67.2 ms]

[Choose CSet: 0.0 ms]

[Ref Proc: 0.5 ms]

[Ref Enq: 0.0 ms]

[Redirty Cards: 62.2 ms]

[Free CSet: 2.8 ms]

[Eden: 21.9G(21.9G)->0.0B(27.2G) Survivors: 8288.0M->2848.0M Heap: 76.0G(94.0G)->57.3G(95.2G)]

[Times: user=53.07 sys=0.05, real=2.37 secs]

Using G1 (-XX:+UseG1GC)

A Minor GC: Stop the World event
O(live set)

2016-02-25T11:58:21.628+0800: [GC pause (G1 Evacuation Pause) (young)

Desired survivor size 8053063680 bytes, new threshold 4 (max 15)

- age 1: 609830392 bytes, 609830392 total
- age 2: 635249376 bytes, 1245079768 total
- age 3: 530928792 bytes, 1776008560 total
- age 4: 6566883776 bytes, 8342892336 total
- age 5: 160917504 bytes, 8503809840 total
, 2.3754150 secs]

Adaptive policy. Threshold goes 4 → 5
looks like most objects don't survive past 4

Many ages may suggest too frequent
collections

[Parallel Time: 2305.9 ms, GC Workers: 23]

[GC Worker Start (ms): Min: 63546061.8, Avg: 63546062.1, Max: 63546062.4, Diff: 0.6]

[Ext Root Scanning (ms): Min: 0.0, Avg: 0.3, Max: 0.6, Diff: 0.5, Sum: 6.3]

[SATB Filtering (ms): Min: 0.0, Avg: 0.0, Max: 0.1, Diff: 0.1, Sum: 0.1]

[Update RS (ms): Min: 16.7, Avg: 18.3, Max: 22.8, Diff: 6.1, Sum: 420.7]

[Processed Buffers: Min: 5, Avg: 10.4, Max: 18, Diff: 13, Sum: 239]

[Scan RS (ms): Min: 247.8, Avg: 252.3, Max: 253.9, Diff: 6.1, Sum: 5803.5]

[Code Root Scanning (ms): Min: 0.0, Avg: 0.0, Max: 0.0, Diff: 0.0, Sum: 0.3]

[Object Copy (ms): Min: 2032.6, Avg: 2033.1, Max: 2034.4, Diff: 1.8, Sum: 46762.1]

[Termination (ms): Min: 0.0, Avg: 1.3, Max: 1.7, Diff: 1.7, Sum: 30.3]

[GC Worker Other (ms): Min: 0.0, Avg: 0.1, Max: 0.2, Diff: 0.2, Sum: 2.1]

[GC Worker Total (ms): Min: 2305.2, Avg: 2305.5, Max: 2305.8, Diff: 0.6, Sum: 53025.4]

[GC Worker End (ms): Min: 63548367.5, Avg: 63548367.6, Max: 63548367.7, Diff: 0.2]

[Code Root Fixup: 0.1 ms]

[Code Root Migration: 0.1 ms]

[Code Root Purge: 0.0 ms]

[Clear CT: 2.1 ms]

[Other: 67.2 ms]

[Choose CSet: 0.0 ms]

[Ref Proc: 0.5 ms]

[Ref Enq: 0.0 ms]

[Redirty Cards: 62.2 ms]

[Free CSet: 2.8 ms]

Adaptive policy: resized generations

[Eden: 21.9G(21.9G)->0.0B(27.2G) Survivors: 8288.0M->2848.0M Heap: 76.0G(94.0G)->57.3G(95.2G)]

[Times: user=53.07 sys=0.05, real=2.37 secs]

2016-02-25T11:58:21.628+0800: [GC pause (G1 Evacuation Pause) (young)

Desired survivor size 8053063680 bytes, new threshold 4 (max 15)

- age 1: 609830392 bytes, 609830392 total
- age 2: 635249376 bytes, 1245079768 total
- age 3: 530928792 bytes, 1776008560 total
- age 4: 6566883776 bytes, 8342892336 total
- age 5: 160917504 bytes, 8503809840 total
, 2.3754150 secs]

[Parallel Time: 2305.9 ms, GC Workers: 23]

[GC Worker Start (ms): Min: 63546061.8, Avg: 63546062.1, Max: 63546062.4, Diff: 0.6]

[Ext Root Scanning (ms): Min: 0.0, Avg: 0.3, Max: 0.6, Diff: 0.5, Sum: 6.3]

[SATB Filtering (ms): Min: 0.0, Avg: 0.0, Max: 0.1, Diff: 0.1, Sum: 0.1]

[Update RS (ms): Min: 16.7, Avg: 18.3, Max: 22.8, Diff: 6.1, Sum: 420.7]

[Processed Buffers: Min: 5, Avg: 10.4, Max: 18, Diff: 13, Sum: 239]

[Scan RS (ms): Min: 247.8, Avg: 252.3, Max: 253.9, Diff: 6.1, Sum: 5803.5]

[Code Root Scanning (ms): Min: 0.0, Avg: 0.0, Max: 0.0, Diff: 0.0, Sum: 0.3]

[Object Copy (ms): Min: 2032.6, Avg: 2033.1, Max: 2034.4, Diff: 1.8, Sum: 46762.1]

[Termination (ms): Min: 0.0, Avg: 1.3, Max: 1.7, Diff: 1.7, Sum: 30.3]

[GC Worker Other (ms): Min: 0.0, Avg: 0.1, Max: 0.2, Diff: 0.2, Sum: 2.1]

[GC Worker Total (ms): Min: 2305.2, Avg: 2305.5, Max: 2305.8, Diff: 0.6, Sum: 53025.4]

[GC Worker End (ms): Min: 63548367.5, Avg: 63548367.6, Max: 63548367.7, Diff: 0.2]

[Code Root Fixup: 0.1 ms]

[Code Root Migration: 0.1 ms]

[Code Root Purge: 0.0 ms]

[Clear CT: 2.1 ms]

[Other: 67.2 ms]

[Choose CSet: 0.0 ms]

[Ref Proc: 0.5 ms]

[Ref Enq: 0.0 ms]

[Redirty Cards: 62.2 ms]

[Free CSet: 2.8 ms]

[Eden: 21.9G(21.9G)->0.0B(27.2G) Survivors: 8288.0M->2848.0M Heap: 76.0G(94.0G)->57.3G(95.2G)]

[Times: user=53.07 sys=0.05, real=2.37 secs]

scanning refs. from other regions

copying objects

~22G out of Eden

~3G survived

Net heap: -19G (3 copied around)

Ouch

2016-02-25T11:14:59.233+0800: [GC pause (G1 Humongous Allocation) (young) (initial-mark)]
Desired survivor size 8053063680 bytes, new threshold 1 (max 15)
- age 1: 9474955328 bytes, 9474955328 total
- age 2: 6322525168 bytes, 15797480496 total
- age 3: 176071416 bytes, 15973551912 total
- age 4: 132526584 bytes, 16106078496 total
, 5.1656688 secs]
[Parallel Time: 5102.7 ms, GC Workers: 23]
[GC Worker Start (ms): Min: 60943668.2, Avg: 60943668.6, Max: 60943668.9, Diff: 0.7]
[Ext Root Scanning (ms): Min: 0.0, Avg: 0.3, Max: 2.5, Diff: 2.5, Sum: 6.2]
[Code Root Marking (ms): Min: 0.0, Avg: 0.3, Max: 3.1, Diff: 3.1, Sum: 6.3]
[Update RS (ms): Min: 13.9, Avg: 17.0, Max: 19.0, Diff: 5.1, Sum: 392.1]
[Processed Buffers: Min: 7, Avg: 10.3, Max: 16, Diff: 9, Sum: 238]
[Scan RS (ms): Min: 301.0, Avg: 302.7, Max: 303.6, Diff: 2.6, Sum: 6962.2]
[Code Root Scanning (ms): Min: 0.0, Avg: 0.0, Max: 0.0, Diff: 0.0, Sum: 0.3]
[Object Copy (ms): Min: 4781.1, Avg: 4781.7, Max: 4782.4, Diff: 1.3, Sum: 109978.1]
[Termination (ms): Min: 0.0, Avg: 0.1, Max: 0.1, Diff: 0.1, Sum: 1.4]
[GC Worker Other (ms): Min: 0.0, Avg: 0.1, Max: 0.2, Diff: 0.1, Sum: 1.9]
[GC Worker Total (ms): Min: 5101.8, Avg: 5102.1, Max: 5102.5, Diff: 0.7, Sum: 117348.6]
[GC Worker End (ms): Min: 60948770.7, Avg: 60948770.7, Max: 60948770.8, Diff: 0.1]
[Code Root Fixup: 0.0 ms]
[Code Root Migration: 0.1 ms]
[Code Root Purge: 0.0 ms]
[Clear CT: 2.7 ms]
[Other: 60.0 ms]
[Choose CSet: 0.0 ms]
[Ref Proc: 1.2 ms]
[Ref Enq: 0.0 ms]
[Redirty Cards: 52.2 ms]
[Free CSet: 4.6 ms]
[Eden: 13.0G(15.0G)->0.0B(23.1G) Survivors: 15.0G->7040.0M Heap: 63.9G(81.4G)->58.0G(86.7G)]
[Times: user=115.69 sys=1.55 real=5.16 secs]

old phase piggy-backing on the young collection

new object > 50% of a region goes straight into OldGen, consume 1 region, and waste part of it

Frequent Humongous? raise -XX: G1HeapRegionSize

copying stuff around again

13G collected
Eden grew

Survivor held
>11G of garbage

Net heap: -6G

```
[Code Root Scanning (ms): Min: 0.0, Avg: 0.0, Max: 0.0, Diff: 0.0, Sum: 0.3]
[Object Copy (ms): Min: 4781.1, Avg: 4781.7, Max: 4782.4, Diff: 1.3, Sum: 109978.1]
[Termination (ms): Min: 0.0, Avg: 0.1, Max: 0.1, Diff: 0.1, Sum: 1.4]
[GC Worker Other (ms): Min: 0.0, Avg: 0.1, Max: 0.2, Diff: 0.1, Sum: 1.9]
[GC Worker Total (ms): Min: 5101.8, Avg: 5102.1, Max: 5102.5, Diff: 0.7, Sum: 117348.6]
[GC Worker End (ms): Min: 60948770.7, Avg: 60948770.7, Max: 60948770.8, Diff: 0.1]
```

STW ended, continues old collection

```
[Code Root Fixup: 0.0 ms]
[Code Root Migration: 0.1 ms]
[Code Root Purge: 0.0 ms]
[Clear CT: 2.7 ms]
[Other: 60.0 ms]
```

Cleanup of 30G not as expensive as previous copying

```
[Choose CSet: 0.0 ms]
[Ref Proc: 1.2 ms]
[Ref Eng: 0.0 ms]
[Redirty Cards: 52.2 ms]
[Free CSet: 4.6 ms]
```

Yet, we were promoting 30G of ephemeral objects

```
[Eden: 13.0G(15.0G)->0.0B(23.1G) Survivors: 15.0G->7040.0M Heap: 63.9G(81.4G)->58.0G(86.7G)]
```

```
[Times: user=115.69 sys=1.55 real=5.16 secs]
```

```
2016-02-25T11:15:04.399+0800: Total time for which application threads were stopped: 5.1662965 seconds
```

```
2016-02-25T11:15:04.399+0800: [GC concurrent-root-region-scan-start]
```

```
2016-02-25T11:15:05.273+0800: [GC concurrent-root-region-scan-end, 0.87349% secs]
```

```
2016-02-25T11:15:05.273+0800: [GC concurrent-mark-start]
```

```
2016-02-25T11:15:06.370+0800: [GC concurrent-mark-reset-for-overflow]
```

```
2016-02-25T11:15:10.371+0800: [GC concurrent-mark-reset-for-overflow]
```

```
2016-02-25T11:15:13.854+0800: [GC concurrent-mark-end, 8.5811484 secs]
```

```
2016-02-25T11:15:13.854+0800: [GC remark [GC ref-proc, 0.0027311 secs], 0.0275626 secs]
```

STW

```
[Times: user=0.43 sys=0.08, real=0.03 secs]
```

```
2016-02-25T11:15:13.882+0800: Total time for which application threads were stopped: 0.0280561 seconds
```

```
2016-02-25T11:15:13.883+0800: [GC cleanup 62G->32G(86G) 0.0453787 secs]
```

STW

```
[Times: user=1.00 sys=0.00, real=0.05 secs]
```

```
2016-02-25T11:15:13.928+0800: Total time for which application threads were stopped: 0.0461448 seconds
```

```
2016-02-25T11:15:13.928+0800: [GC concurrent-cleanup-start]
```


```
2016-02-25T11:15:13.943+0800: [GC concurrent-cleanup-end, 0.0148144 secs]
```



Memory utilisation

App not responsive for > 2s, >6s. What is going on?

- Lots of churn (objects created, soon dereferenced)
 - `jstat`, `jcmd` help seeing at what rate (or Flight Recorder if accessible)
- Allocation pressure → frequent Young GC
 - More copies between survivor spaces
 - Premature promotions create *collection debt*
- Humongous allocations
 - fragmentation → inefficient memory utilisation → even more GC work
- Clients unhappy and/or services downstream cascading failures
 - <https://www.elastic.co/blog/elastic-cloud-outage-april-2016>
ZooKeeper (coordination service) dies for GC, major outages in Elastic cloud

GCViewer: gc log visualization





GC tuning

Worth several talks in itself. Some knobs that might be relevant..

- `-XX:MaxGCPauseMillis=200` Informs adaptive policies
- `-XX:+PrintReferenceGC` Details into object references
- `-XX:+AlwaysTenure` Straight to Old Gen (spare copies)
- `-XX:+NeverTenure` Never to Old Gen (assume mostly garbage)
- `-XX:+BindGCTaskThreadsToCPUs`
- `-XX:CMSInitiatingOccupancyFraction`
- `-XX:InitiatingHeapOccupancyPercent` **Tolerance to utilisation**
- `-XX:+ScavengeBeforeFullGC` / `-XX:+CMSScavengeBeforeRemark`

Collect eden before Full GC or CMS Remark, sparing the cross-generation ref checks

Application tuning

“The demand upon a resource tends to expand to match its supply.”

~ Parkinson’s Law

Adding memory generally delays, not fixes, problems.


More space → more garbage → more copies → more collections

Complementary approach

- What (ab)uses these resources? How? Why?


Misbehaviour example: GC nepotism

A Linked list




Minor GC

Two objects removed..






Minor GC

... and collected





All good




Misbehaviour example: GC nepotism



Minor GC



Minor GC



Minor GC

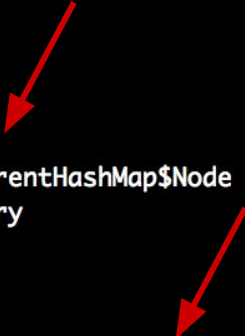
Minor GC doesn't clean old Gen: the
deleted object holds a reference. The
2nd survives until promotion.

Measuring memory footprint

```
$ jmap -heap $PID
$ jmap -histo $PID
$ jmap -histo:live $PID
```

Compare with `-histo:live` (full GC, so do it last)

num	#instances	#bytes	class name
1:	19042	9349608	[B
2:	32861	3747800	[C
3:	3142	2107224	[I
4:	20607	494568	java.lang.String
5:	3212	354976	java.lang.Class
6:	4376	291776	[Ljava.lang.Object;
7:	2586	227568	java.lang.reflect.Method
8:	1256	188752	[Ljava.util.HashMap\$Node;
9:	5202	166464	java.util.concurrent.ConcurrentHashMap\$Node
10:	3326	133040	java.util.LinkedHashMap\$Entry
11:	3294	105408	java.util.HashMap\$Node
12:	2346	93840	java.lang.ref.Finalizer
13:	1259	90648	java.lang.reflect.Field
14:	3270	79096	[Ljava.lang.Class;
15:	59	76112	[Ljava.util.concurrent.ConcurrentHashMap\$Node;



[Z = boolean

[B = byte

[S = short

[I = int

[J = long

[F = float

[D = double

[C = char

[L = any non-primitives(Object)

Examining heap

Runtime dump

```
jmap -dump:format=b,file=dump.hprof $PID  
jhat dump.hprof  
[...]
```

Started HTTP server on port 7000

JVM dump on error

```
-XX:HeapDumpPath=/var/log/my-service/  
-XX:+HeapDumpOnOutOfMemoryError  
-XX:+HeapDumpAfterFullGC  
-XX:+HeapDumpBeforeFullGC
```

Also: Eclipse MAT, Visual VM, Mission Control...

- Lots of candy: track dominator trees, map collisions, object ages, OQL
- Mission control allows defining triggers based on behaviour

Object graph: <http://openjdk.java.net/projects/code-tools/jol/>

Consider the cost of abstractions

num	#instances	#bytes	class name
1:	456735295	29968183048	[C
2:	141993549	17650832184	[Ljava.lang.Object;
3:	432874195	13851974240	java.lang.String
4:	141783960	5671358400	java.util.ArrayList
5:	220867901	5300829624	java.lang.Long
6:	3507261	3992725000	[I
7:	90242360	3839836936	[B
8:	142089373	3410144952	.JDBCRecord

num	#instances	#bytes	class name
1:	309717286	18357971096	[C
2:	103220845	12801064160	[Ljava.lang.Object;
3:	309717154	9910948928	java.lang.String
4:	103219696	4128787840	java.util.ArrayList
5:	103216209	2477189016	.JDBCRecord
6:	103005153	2472123672	java.lang.Long
7:	5741	22010304	[B
8:	211348	5072352	java.lang.Double

Boxing:

- 5.3G / 220M Long instances = 24 bytes
- Boxing: 3x overhead on long (8 bytes)

Scala (closures, Java conversions, immutables)

Strings:

```
String copy = new String(a + b) // NO
String copy = a + b             // YES
```


-XX:+PrintStringTableStatistics

-XX:+UseStringDeduplication (only G1)

Consider the cost of abstractions

Object headers

ordinary
object pointer



<http://hg.openjdk.java.net/jdk8/jdk8/hotspot/file/tip/src/share/vm/oops/oop.hpp>

<http://hg.openjdk.java.net/jdk8/jdk8/hotspot/file/tip/src/share/vm/oops/markOop.hpp>

- 64-bit: 12 bytes padded to multiple of 8 → 16 bytes
- 32-bit: 8 bytes padded to multiple of 4 → 12 bytes

References

- Ref = 4 bytes on < 32G heaps
- Ref = 8 bytes on 64-bit JVMs with >32G heaps

Arrays: 1 ref to type, 4 bytes for length, 1 ref per element. Min 8/16 bytes

Consider the cost of abstractions

- Boxing

```
class A {  
    byte x;  
}  
class B extends A {  
    byte y;  
}  
  
new A() = 24 bytes  
new B() = 32 bytes  
class C {  
    Object o = new Object();  
}  
new C() = 40 bytes..
```

- Growing heap from 24G -> 48G? Think crossing tax brackets..
- -XX:+UseCompressedOops will compresses native pointer to 32bits
 - <https://wiki.openjdk.java.net/display/HotSpot/CompressedOops>
 - Should be enabled in recent JVMs

Consider the cost of abstractions

```
while ((line = reader.readLine()) != null) {  
    users.add(new User(line));  
}
```

```
class User {  
    private final String name;  
    private final Date birth;  
    ...  
    public User(String s) {  
        String[] fields = s.split("::");  
        this.name = fields[0];  
        this.birth = dateFormat.parse(fields[1]) ;  
        ...  
    }  
    public String getName() { .. }  
    public Date getBirth() { return new Date(birth.getTime) }  
    ...  
    public String getXXX()  
}
```

Good OOP, trying to save CPU on access.. but..

Can we afford multiplying dataset sizes?

Does our internal representation need to mirror the public contract?

Consider the cost of abstractions

```
while ((line = reader.readLine()) != null) {  
    users.add(new User(line));  
}
```

```
class User {  
    private final String data;  
  
    ...  
  
    public User(String s) {  
        this.data = s;  
    }  
  
    public String getName() {  
        return findField(0);  
    }  
  
    public Date getBirth() {  
        return new Date(findField(1))  
    }  
  
    ...  
  
    private String findField(int n) {  
        // loop to find field  
    }  
}
```

← Might make sense.. (or, store offsets but not parse) to delay allocation until it's really needed

- Trades CPU (hardly saturated) for memory
- Think more complex cases:

```
class Ethernet implements L2 {  
    MAC src; MAC dst; Short[] vlans;  
    L3Packet payload;  
}
```

```
class IPPacket implements L3 {  
    IP src; IP dst; Flags flags;  
    L4Datagram payload;  
}
```

Unaffordable with millions of instances..

Calculation of object size (only Hotspot)

```
import jdk.nashorn.internal.ir.debug.ObjectSizeCalculator;
import static jdk.nashorn.internal.ir.debug.ObjectSizeCalculator.*

ObjectSizeCalculator sizeCalc = new ObjectSizeCalculator(
    getEffectiveMemoryLayoutSpecification());

long size = sizeCalc.getObjectSize(new Record(...));
```

Consider the cost of abstractions

Data from/to disk/network into objects implies going through multiple copies..

Typical case:

disk/network → kernel buffer → userspace buffer → byte[] → String → Objects




Memory mapped files (useful for large files, IPC..)

```
ByteBuffer b = fileChannel.map(READ_ONLY, 0, file.size())
```

Direct memory buffers (self-managed memory)

```
ByteBuffer dbb = ByteBuffer.allocateDirect(file.size())  
fileChannel.read(directByteBuffer)
```

Consider the cost of abstractions



```
ByteBuffer data = ...;
...
while (data.hasRemaining()) {
    ByteBuffer bytes = data.slice()
    bytes.limit(RECORD_SIZE)
    data.position(data.position() + RECORD_SIZE)
    users.add(new Record(bytes))
}
```

- Easy to build an `Iterator[Record]` over a `ByteBuffer`
- Single copy of the data
- Easier to achieve cache friendliness

Examining off-heap memory

Tracking native allocations

JVM flag required

```
-XX:NativeMemoryTracking=off | summary | detail
```

Retrieve info

```
$ jcmd $PID VM.native_memory baseline           # set  
$ jcmd $PID VM.native_memory summary.diff      # poll for diff
```

Unsuspected memory sinks

Unsuspected memory sinks lurk everywhere... know your APIs, libraries

- **Logs** ~ `log.debug("Request " + req.id + " is generating useless garbage")`
 - Strings, Message objects, locks ...
- **Lazy Initialization in Scala**: additional int, + sync overhead
- `ArrayList.addAll` → allocates an `Object[size]`
- An object with a `finalize()` method allocates an additional object
 - You don't want `finalize()` on classes with millions of instances
 - Takes 2 GC cycles to clean
- `WeakHashMap` has a delay to clean dead refs (lazy eviction)
- **Secret NIO `ByteBuffer` cache** avoids expensive `malloc` / `free` sequences for short lived buffers... by potentially caching massive buffers
 - <http://mail.openjdk.java.net/pipermail/nio-dev/2015-December/003420.html>

CPU

JIT optimisations: Escape analysis + Inlining

```
public A {  
    private final int x;  
    public A(final int _x) {  
        this.x = _x;  
        this.y = _y;  
    }  
    public int getY() { return y; }  
}
```

```
public void f(int n) {  
    int x = 0;  
    for (i = 0; i < n; i++) {  
        A a = new A(i);  
        System.out.println(a.getX());  
    }  
}
```

(likely) JIT'ed version

```
public void f(int n) {  
    int x = 0;  
    for (i = 0; i < n; i++) {  
        int _x = x;  
        System.out.println(_x);  
    }  
}
```

Objects that don't escape curr. method or thread might get stack allocation.

Methods calls may get inlined

Consider the cost of abstractions

- JIT vs OOP: Megamorphic methods can't be optimized

<https://github.com/google/guava/issues/1268>

"... guava Immutable collections [...] have specializations for zero (EmptyImmutableList) and one (SingletonImmutableList) element collections. These specializations take the form of subclasses of ImmutableList, to go along with the "Regular" implementation and a few other specializations like ReverseImmutable, SubList, etc.

The result is that when these subclasses mix at some call site, the call is megamorphic, and performance is awful compared to classes without these specializations (worse by a factor of 20 or more)."

JIT: Escape analysis + Inlining

Very relevant for Scala, Java8 lambdas

```
def maybeDouble(Option[Int] o): Option[Long] = {  
    o.map { _ * 2 } // o.map(new Function(x: Int) { return x * 2; })  
}
```

Help the JIT help you

- Small functions, clean code, immutability, few conditionals, avoid megamorphism
- Profile allocations & benchmark performance to validate assumptions
- JIT watch: <https://github.com/AdoptOpenJDK/jitwatch>
- Gil Tene: <http://infoq.com/presentations/java-jit-optimization>

Latency jitter & spikes

```
2016-05-10T17:06:19.340+0800: Total time for which application threads were stopped: 0.0010981 seconds
2016-05-10T17:06:19.341+0800: Total time for which application threads were stopped: 0.0009505 seconds
2016-05-10T17:06:19.342+0800: Total time for which application threads were stopped: 0.0008453 seconds
2016-05-10T17:06:19.343+0800: Total time for which application threads were stopped: 0.0008495 seconds
```

These are not necessarily due to GC. More info using:

```
-XX:+UnlockDiagnosticVMOptions -XX:+PrintSafepointStatistics
```

Time to safepoint. Identified by: 12.754: no vm operation

- Some JVMs introduce a periodic guaranteed safepoint time (used to perform GC and other tasks, e.g.: apply/revoke code optimisations)
- Can be controlled with `-XX:GuaranteedSafepointInterval=300000`
- <http://epickrram.blogspot.com.es/2015/08/jvm-guaranteed-safepoints.html>

Latency jitter & spikes

```
2016-05-10T17:06:19.340+0800: Total time for which application threads were stopped: 0.0010981 seconds
2016-05-10T17:06:19.341+0800: Total time for which application threads were stopped: 0.0009505 seconds
2016-05-10T17:06:19.342+0800: Total time for which application threads were stopped: 0.0008453 seconds
2016-05-10T17:06:19.343+0800: Total time for which application threads were stopped: 0.0008495 seconds
```

These are not necessarily due to GC. More info using:

```
-XX:+UnlockDiagnosticVMOptions -XX:+PrintSafepointStatistics
```

Biased locking. Identified by: 26.319: RevokeBias..

Optimizes contended locks: last accessor thread has higher chances on next attempt Pro: cache friendliness ; Con: bookkeeping, bad on thread pools, highly concurrent apps..

- Disable with `-XX:-UseBiasedLocking`

Stack dumps

```
$ jstack -l $PID > output
```

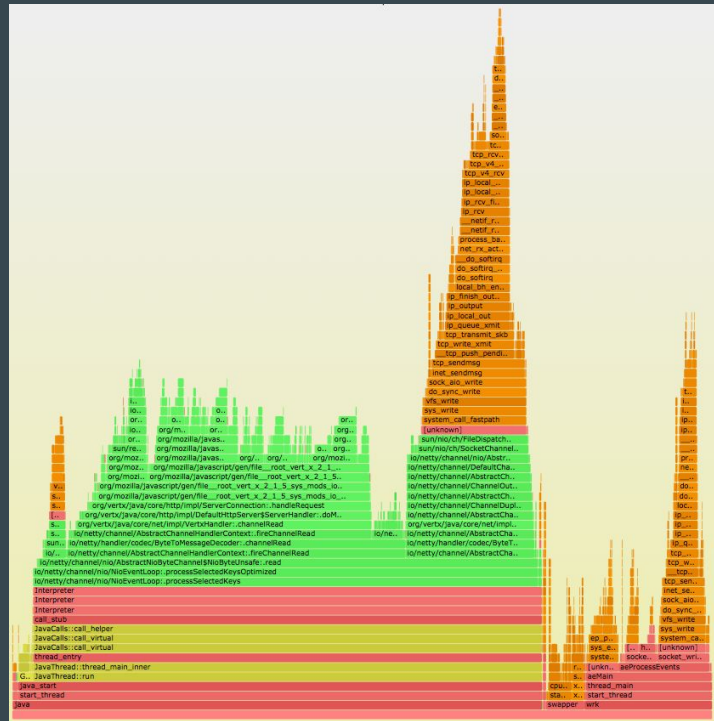
```
$ kill -3 $PID # goes to stderr, wherever this is directed to
```

- -l gives additional info about locks
- <https://github.com/spotify/threaddump-analyzer>

WARN!

- A thread's stack is only retrieved at safepoints (most JVMs)
 - Hurts accuracy of reported stacks
 - This also affects profilers
- One thread's stack dumped at a time
 - Inconsistent stacks: two threads hold the same lock; thread blocked on free monitor.. (you can see this in a few slides)

FLAME GRAPHS



Stack visualization, crossing JVM → OS

Very effective to spot where CPU time is going

Drill down to specific sections of the stack

Synchronization

```
"pool-1-thread-55" #657 prio=5 os_prio=0 tid=0x00007f861702d000 nid=0xc4cc runnable [0x00007f8611bdc000]
  java.lang.Thread.State: RUNNABLE
    at sun.nio.ch.FileDispatcherImpl.read0(Native Method)
    at sun.nio.ch.SocketDispatcher.read(Unknown Source)
    at sun.nio.ch.IOUtil.readIntoNativeBuffer(Unknown Source)
    at sun.nio.ch.IOUtil.read(Unknown Source)
    at sun.nio.ch.SocketChannelImpl.read(Unknown Source)
    - locked <0x00007f894c1d0158> (a java.lang.Object)
    at .serializer.MessageBuffer.readPartial(MessageBuffer.java:211)
    - locked <0x00007f8957ef3d38> (a .serializer.MessageBuffer)
```

native thread ID
(find with top / ps / htop..)

Two monitors held while we don't do anything (waiting for OS on an IO read)

Synchronization

```
"pool-1-thread-40" #82 prio=5 os_prio=0 tid=0x00007f8614367000 nid=0xa955 waiting for monitor entry [0x00007f861cccc000]
java.lang.Thread.State: BLOCKED (on object monitor)
    at      .log.CompressFileHandler.publish(CompressFileHandler.java:487)
    - waiting to lock <0x00007f87463d4930> (a      .log.CompressFileHandler)
    at java.util.logging.Logger.log(Unknown Source)
    at java.util.logging.Logger.doLog(Unknown Source)
    at java.util.logging.Logger.log(Unknown Source)
    at java.util.logging.Logger.warning(Unknown Source)
    at      .server.handler.ProxyConnectionTransaction.startTransaction(ProxyConnectionTransaction.java:123)
```

```
"pool-1-thread-38" #80 prio=5 os_prio=0 tid=0x00007f861431f800 nid=0xa951 runnable [0x00007f861cece000]
```

```
java.lang.Thread.State: RUNNABLE
```

```
at java.util.logging.StreamHandler.flush(Unknown Source)
```

```
- locked <0x00007f87463d4930> (a .log.CompressFileHandler)
```

```
at .log.CompressFileHandler.publish(CompressFileHandler.java:491)
```

```
- locked <0x00007f87463d4930> (a .log.CompressFileHandler)
```

```
at java.util.logging.Logger.log(Unknown Source)
```

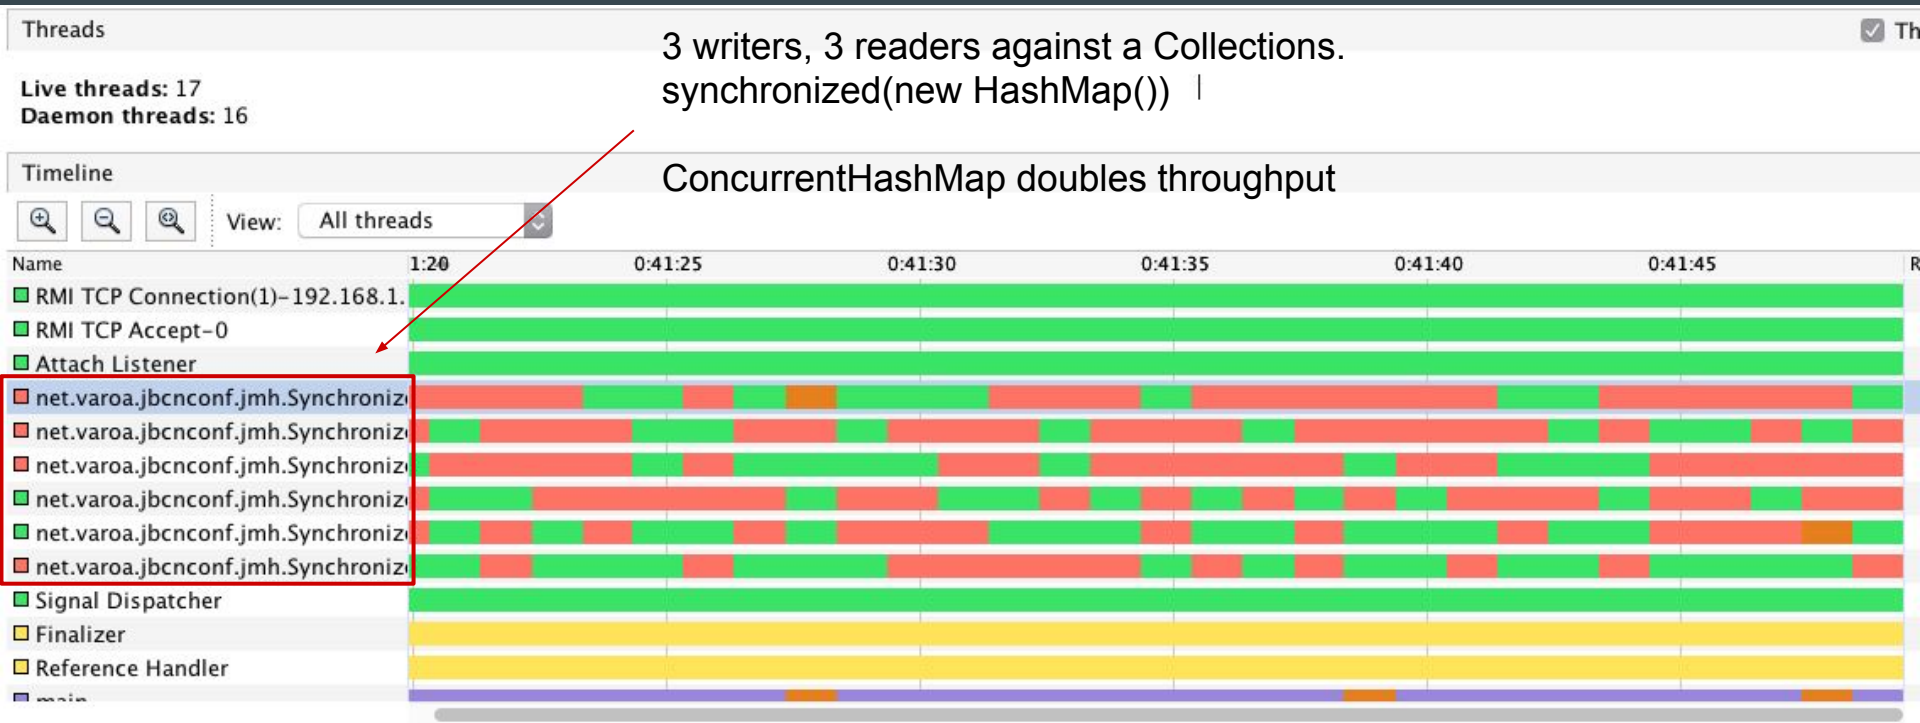
```
at java.util.logging.Logger.doLog(Unknown Source)
```

```
at java.util.logging.Logger.log(Unknown Source)
```

```
at java.util.logging.Logger.info(Unknown Source)
```

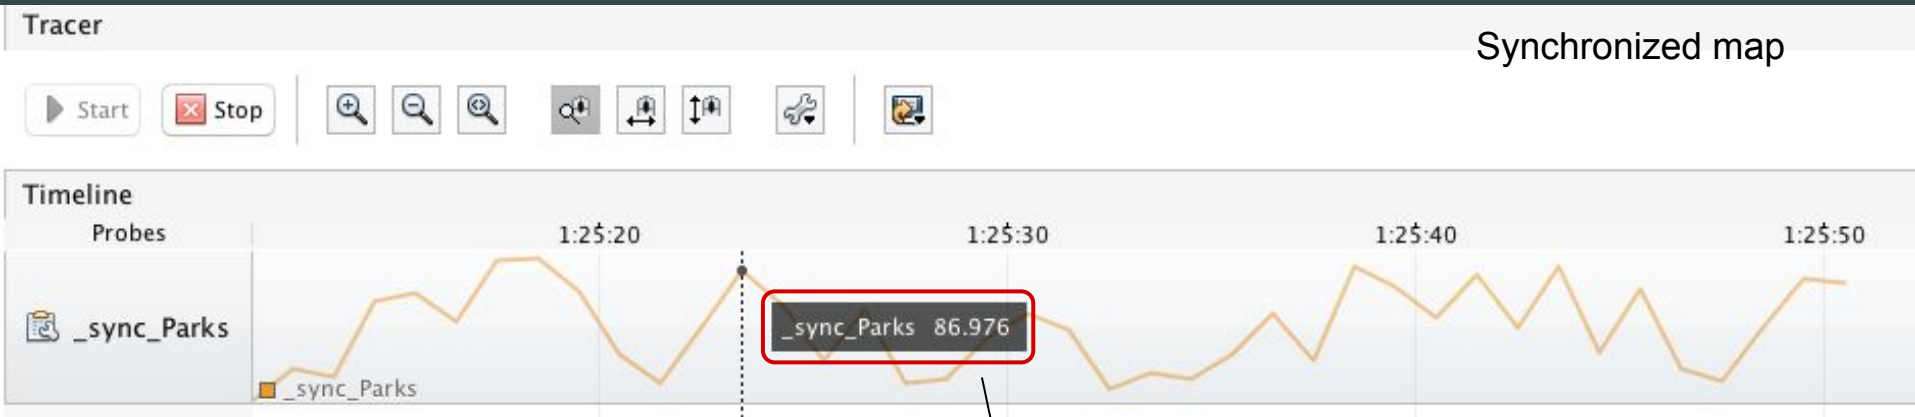
Lock held during IO,
blocking *anyone* trying
to log a message

Synchronization



Synchronization

```
jcmd $PID PerfCounter.print | grep Parks
```



Synchronization

Synchronized blocks imply serial parts of the program

Amdahl's law:
$$S_{\text{latency}}(s) = \frac{1}{(1 - p) + \frac{p}{s}}$$

- A x10 speedup of 10% of the exec. time ($p = 0.1, s = 10$) $\rightarrow \sim 1.10$ speedup
- A x1.5 speedup of 90% of the exec. time ($p = 0.9, s = 1.5$) $\rightarrow \sim 1.43$ speedup

Avoid

- Large synchronized blocks
- Synchronizing on `this` (you're becoming vulnerable to potential blocks)
- Calling external methods while synchronized (e.g.: see below, in the JDK)

```
public synchronized String formatMessage(LogRecord record) {
```

CPU utilisation (or lack of)

```
perf stat -d -p $PID      # also: cat /proc/$PID/status
```

Performance counter stats for process id '46185':

46802,773132	task-clock	#	0,367 CPUs utilized	
24521	context-switches	#	0,001 M/sec	
1056	CPU-migrations	#	0,000 M/sec	
450	page-faults	#	0,000 M/sec	
129724461306	cycles	#	2,772 GHz	[40,03%]
120451823606	stalled-cycles-frontend	#	92,85% frontend cycles idle	[40,40%]
101327998695	stalled-cycles-backend	#	78,11% backend cycles idle	[40,56%]
5710255927	instructions	#	0,04 insns per cycle	
		#	21,09 stalled cycles per insn	[50,69%]
1334448420	branches	#	28,512 M/sec	[50,80%]
34780770	branch-misses	#	2,61% of all branches	[50,69%]
1540840733	L1-dcache-loads	#	32,922 M/sec	[50,22%]
4551850342	L1-dcache-load-misses	#	295,41% of all L1-dcache hits	[50,02%]
71937736	LLC-loads	#	1,537 M/sec	[39,83%]
42543914	LLC-load-misses	#	59,14% of all LL-cache hits	[39,78%]

127,428135422 seconds time elapsed

Concurrency toolbox

The JDK offers resources for concurrency

- `java.util.concurrent.locks.*`
- `java.util.concurrent.*` → e.g.: thread safe collections

Does your application really need locks?

- `volatile` : writes guaranteed to be visible when other threads read
- **AtomicXX** classes, operations: `incrementAndGet`, `compareAndSwap ...`
- `AtomicXX.lazySet()` : writes not reordered with later writes (cheaper for single writer)

High performance lock-free collections: <https://github.com/JCTools/JCTools>

Concurrency toolbox: coordination

Example:

```
Queue<T> queue; // unbounded queue, we can't change the implementation
...
public void onNext(T t) {
    queue.offer(t) // enqueues a new item
    pollQueue(wip, requested, queue, child) // drains `requested` items from the queue
}
```

How to bound the queue, and trigger a backpressure notification?

```
synchronized(queue) {
    if (queue.size() < limit)
        queue.offer(t);
    else
        callback.apply()
        return false;
}
```

← serial block and risk of parking threads

← what if this blocks? we're effectively blocking anyone else accessing the queue, even consumers

AtomicInteger

Snapshot state

If full, make
sure exactly 1
thread deals
with the
consequences

Confirm slot for
our item if
nothing changed

```
94  @Override
95  +    public void onNext(T t) {
96  +        if (!ensureCapacity()) {
97  +            return;
98  +        }
99      queue.offer(on.next(t));
100  +    pollQueue(wip, requested, capacity, queue, child);
101  }
102
103  +    private boolean ensureCapacity() {
104  +        if (capacity == null) {
105  +            return true;
106  +        }
107  +
108  +        long currCapacity;
109  +        do {
110  +            currCapacity = capacity.get();
111  +            if (currCapacity <= 0) {
112  +                if (saturated.compareAndSet(false, true)) {
113  +                    // ensure single completion contract
114  +                    child.onError(new BufferOverflowException());
115  +                    unsubscribe();
116  +                    if (onOverflow != null) {
117  +                        onOverflow.call();
118  +                    }
119  +                }
120  +                return false;
121  +            }
122  +            // ensure no other thread stole our slot, or retry
123  +            } while (!capacity.compareAndSet(currCapacity, currCapacity - 1));
124  +            return true;
125  +        }
126  }
```

If contended, threads collaborate to
perform actions, rather than block
each other

Concurrency toolbox: ThreadLocal

```
class Formatter {
    DateFormat df = new SimpleDateFormat("DDMMYYYY")
    public String format(Date d) {
        return df.format(d);
    }
}

// SDF is not thread safe ..
// .. so this is neither
// Can we avoid creating an
// instance per .format() call?

class Formatter {
    ThreadLocal<SimpleDateFormat> df =new ThreadLocal() { // Wrap it in ThreadLocal
        public SimpleDateFormat initialValue() {           // Called on the first .get()
            return new SimpleDateFormat();                  // performed by each Thread
        }
    };
    public String format(Date d) {
        return df.get().format(d);
    }
}

// Now thread-safe as each
// thread gets its own instance
// of SimpleDateFormat
```


Concurrency toolbox: Thread Local buffers

This is actually how the JVM deals with allocations from multiple threads. Avoids contention by allocating on Thread-Local Allocation Buffers.

- Enabled by default: `-XX:+UseTLAB`
- See details in Flight Recorder (next slide), or: `-XX:+PrintTLAB`

```
TLAB: gc thread: 0x00007f3c1ff0f800 [id: 10519] desired_size: 221KB
      slow allocs: 8   refill waste: 3536B alloc: 0.01613      11058KB
      refills: 73 waste 0.1% gc: 10368B slow: 2112B fast: 0B
```

- `-XX:+ResizeTLAB` let the JVM dynamically adjust size
- `-XX:TLABSize=2m -XX:MinTLABSize=64k` adjust it yourself
- `-XX:+AggressiveOpts`

Allocations

Events Operative

Interval: 52 s 952 ms (all)

☐ Synchronize Selection

22/05/16 20:51:37

22/05/16 20:52:30

General Allocation in New TLAB Allocation Outside TLABs

Thread Local Allocation Buffer (TLAB) Statistics


TLAB Count	11.024
Maximum TLAB Size	2,01 MB
Minimum TLAB Size	2,03 kB
Average TLAB Size	722,39 kB
Total Memory Allocated for TLABs	7,59 GB
Allocation Rate for TLABs	146,87 MB/s



Statistics for Object Allocations (Outside TLABs)

Object Count	647
Maximum Object Size	2,88 kB
Minimum Object Size	16 bytes
Average Object Size	830 bytes
Total Memory Allocated for Objects	524,91 kB
Allocation Rate for Objects	9,91 kB/s

Allocation

☒ TLAB Allocations ☒ Object Allocations (Outside TLABs)

Allocations

Events Operative

Interval: 2 min 15 s (all)

☐ Synchronize Selection


22/05/16 23:59:17

23/05/16 0:01:33

General Allocation in New TLAB Allocation Outside TLABs

Allocation by Class Allocation by Thread Allocation Profile

Allocation Pressure



Thread	Average Object Size	TLABs	Total TLAB Size	Pressure
qtp1512981843-15	1,010 bytes	356	36,05 MB	13,35%
qtp1512981843-18	1,01 kB	353	35,15 MB	13,02%
qtp1512981843-19	966 bytes	351	34,57 MB	12,80%
qtp1512981843-16	1,01 kB	348	34,38 MB	12,73%
qtp1512981843-17	988 bytes	349	34,29 MB	12,70%
qtp1512981843-13	921 bytes	350	34,15 MB	12,65%
qtp1512981843-12	972 bytes	336	34,06 MB	12,61%
qtp1512981843-14-acceptor-0@518b4cf1-ServerConnector@5f4...	47 bytes	255	12,10 MB	4,48%
runner	40 bytes	239	10,04 MB	3,72%
Service Thread	39 bytes	19	1,61 MB	0,60%
RMI TCP Connection(7)-192.168.1.134	112 bytes	19	840,16 kB	0,30%
RMI TCP Connection(6)-192.168.1.134	259 bytes	23	699,59 kB	0,25%
Scheduler-1285044316	32 bytes	38	669,40 kB	0,24%
C1 CompilerThread2	24 bytes	8	386,16 kB	0,14%
RMI TCP Connection(idle)	24 bytes	2	295,29 kB	0,11%

Stack Trace

Stack Trace	TLABs	Total TLAB Size	Pressure

Unintended contention

False sharing

```
class MyClass[T] {  
    public volatile long a = 0;  
    public volatile long b = 0;  
}
```

Thread A writes to a in CPU1

Thread B reads b in CPU2

Do we have contention? Yes

64 bytes

object header	a(8)	b(8)	other things
---------------	------	------	--------------

The CPU works with a full cache line, not individual fields.

A takes exclusive ownership of the full cache line, updates and B's copy is invalidated, even if B's value didn't change. Same applies to the rest of the line

Unintended contention

False sharing

```
class MyClass[T] {  
    PaddedAtomicLong a;  
    PaddedAtomicLong b;  
}
```

```
class PaddedAtomicLong extends AtomicLong {  
    public final long  
        p1, p2, p3, p4, p5, p6 = 7L;  
}
```

64 bytes

object header	*a	*b	
object header	val		
object header	val		

Padding ensures they reside on
different cache lines

<http://mechanical-sympathy.blogspot.com.es/2011/07/false-sharing.html>

<http://mechanical-sympathy.blogspot.com.es/2011/08/false-sharing-java-7.html>

<http://psy-lob-saw.blogspot.com.es/2014/03/java-object-layout-tale-of-confusion.html>

Leftovers

- Java Microbenchmark Harness (JMH)
- Have performance targets as business requirements
- Perf tests are hard (e.g.: generating more load than production)
- Instrumentation, monitorization
 - As part of CI: add performance tests early, check for regressions...
 - Production: make behaviour visible, spot anomalies
- No need to optimize early, but have a story for how you'd improve when needed

Q & A

Thanks!